AWS Cloud Labs

Complete Lab Guide — Weeks 1 to 11

Week 1 – EC2 Instance: Linux & Windows Server

Pre-Lab Questions & Answers

Q1. What is EC2?
EC2 stands for Amazon Elastic Compute Cloud. It is a web service that provides secure, scalable compute capacity in the cloud. It is a virtual server in Amazon's infrastructure for running applications on AWS.
Q2. What is an Instance?
An instance is a virtual server in the AWS cloud. It runs on EC2 and can be configured with different operating systems, CPU, memory, and storage based on the chosen instance type (e.g., t2.micro).
Q3. What is the purpose of a Security Group?
A security group acts as a virtual firewall for EC2 instances. It controls incoming (inbound) and outgoing (outbound) traffic. Inbound rules control what traffic can reach your instance; outbound rules control what traffic leaves it.
Q4. What is a Key Pair?
A key pair consists of a public key and a private key used as security credentials to connect to an EC2 instance. AWS stores the public key on the instance; you download and keep the private key (.pem file). For Linux, it is used with SSH/PuTTY; for Windows, it is used to decrypt the RDP password.
Q5. List the Operating Systems supported by EC2.
Amazon Linux, Ubuntu, Windows Server, Red Hat Enterprise Linux (RHEL), SUSE Linux Enterprise Server, openSUSE Leap, Fedora, Fedora CoreOS, Debian, CentOS, Gentoo Linux, Oracle Linux, FreeBSD, and macOS.

Step-by-Step: Launch a Linux Server (EC2)

A. Create the EC2 Instance

- Go to the AWS Console → Select EC2 Service
- Click on Launch Instances
- Enter a Name for the instance

- Choose Amazon Machine Image (AMI) → select Amazon Linux
- Choose Instance Type → select t2.micro
- Click on Create New Key Pair
  - Enter a Key Pair Name
  - Key Pair Type: RSA
  - Private Key File Format: .pem
  - Click Create Key Pair → save the .pem file to your computer
- Under Network Settings → Enable SSH Port (Port 22)
- Configure Storage as needed
- Click Launch Instance

B. Connect via PuTTY (SSH Client Method)

- Download and install PuTTY and PuTTYgen from: https://www.chiark.greenend.org.uk/~sgtatham/putty/latest.html
- Convert .pem to .ppk using PuTTYgen:
  - Open PuTTYgen → Load → select your .pem file → Save Private Key → save as .ppk
- In the EC2 Console → select your instance → click Connect → go to SSH Client tab → note the Hostname (e.g., ec2-xx-xx-xx-xx.compute.amazonaws.com)
- Open PuTTY:
  - Enter the Hostname
  - Go to Connection → SSH → Auth → Credentials → Browse → select the .ppk file
  - Click Open → Accept the security alert → login as: ec2-user

C. Connect via EC2 Instance Connect (Browser Method)

- Select the instance → click Connect → go to EC2 Instance Connect tab
- Note the Public IPv4 Address
- Open PuTTY → Enter the IP Address → attach the .ppk file → click Open → enter username: ec2-user

---

Step-by-Step: Launch a Windows Server (EC2)

- Select the AWS Region
- Go to EC2 Service → click Launch Instances
- Enter a Name for the instance
- Choose AMI → select Windows Server
- Choose Instance Type → t2.micro
- Click Create New Key Pair:
  - Enter a Key Pair Name (no special characters)
  - Private Key File Format: .pem
  - Click Create Key Pair
- Under Network Settings → check Allow RDP Traffic → Source: Anywhere

- Configure Storage → Click Launch Instance
- To stop or remove: Select the instance → click Stop or Terminate

Week 2 – Transfer Files: Local System to Linux & Windows Server

Pre-Lab Questions & Answers

Q1. Define Server.
A server is a computer or software system that provides resources, data, services, or programs to other computers (clients) over a network. In AWS, an EC2 instance acts as a virtual server.
Q2. What is an Instance?
An instance is a virtual server in EC2. It can be configured with different operating systems and hardware specifications to run your applications in the cloud.
Q3. What is the purpose of a Security Group?
A security group is a virtual firewall that controls inbound and outbound traffic to your EC2 instance. It defines which ports and protocols are open or blocked.
Q4. What is a Key Pair?
A key pair (public key + private key) is used to securely authenticate to EC2 instances. For Linux, it enables SSH login; for Windows, it is used to decrypt the administrator password for RDP access.
Q5. List the Operating Systems supported by EC2.
Amazon Linux, Ubuntu, Windows Server, RHEL, SUSE Linux, openSUSE Leap, Fedora, Fedora CoreOS, Debian, CentOS, Gentoo Linux, Oracle Linux, FreeBSD, and macOS.

Step-by-Step: Transfer Files to Windows Server

- Create a Windows EC2 Machine (refer to Week 1 steps)
- Generate the Password:
  - Select instance → Actions → Security → Get Windows Password
  - Upload the .pem key → Decrypt Password → note the password
- Create a file on your Local Machine that you want to transfer
- Open Remote Desktop Connection (RDP) on your local PC:
  - Click Show Options
  - Go to the Local Resources tab
  - Click More under Local Devices and Resources
  - Check the Drives checkbox (this shares your local drives) → click OK
- Click Connect → enter the EC2 Public IP, username (Administrator), and password
- Once connected, open File Explorer inside the EC2 machine → your local drive is visible
- Copy the file from your local drive and paste it into the EC2 machine

Step-by-Step: Transfer Files to Linux Server

- Create a Linux EC2 Machine (refer to Week 1 steps)
- Generate the .ppk file using PuTTYgen (refer to Week 1 steps)
- Download and Install WinSCP from: https://winscp.net/eng/download.php
- Open WinSCP:
  - Enter the Host Name (EC2 Public IP)
  - Enter the Username: ec2-user
  - Click Advanced → SSH → Authentication → attach your .ppk file → click OK
- Click Login → click Yes on the security alert
- Drag and drop the file from your local machine into the Linux server pane in WinSCP
- To verify: Open PuTTY → connect to the Linux server → use ls command to check the transferred files

Week 3 – Configure Linux & Windows as Web Servers

Pre-Lab Questions & Answers

Q1. What is a Web Server?
A web server is a program that uses HTTP (Hypertext Transfer Protocol) to serve web page files to users in response to their requests. Examples include Apache HTTP Server (HTTPD), Internet Information Services (IIS), and NGINX.
Q2. What is an Instance?
An instance is a virtual server running on EC2 in the AWS cloud, configured with an OS and resources to run applications.
Q3. What is a Repository?
A Linux repository is a storage location from which the operating system retrieves and installs software updates and applications.
Q4. What is a YUM Repository?
YUM (Yellowdog Updater Modified) Repositories are warehouses of Linux software packages. When software is not available by default in Linux, we use YUM repositories to find and install it. Command: yum install <packagename>
Q5. How to Create Empty Files in Linux?
Use the touch command: sudo touch filename — This creates an empty file with the given name. Example: sudo touch file1 file2 file3

Step-by-Step: Linux Web Server (HTTPD/Apache)

HTTPD (HTTP Daemon) is a background software program that acts as a web server. It listens for incoming requests and serves HTML/multimedia documents over the internet using HTTP.

Step 1 – Connect and Switch to Root

- Create and connect to a Linux EC2 Machine (refer to Week 1)
- Switch to root user:

```
sudo su
```

Step 2 – Update YUM Repository

```
yum update -y
```

Step 3 – Install HTTPD Web Package

```
yum install httpd -y
```

Step 4 – Create a Web Page

- Navigate to the web directory:

```
cd /var/www/html
```

- Create the index page:

```
echo "Amazon Web Services" > index.html
```

Step 5 – Start the Web Server

```
service httpd start
chkconfig httpd on
```

Step 6 – Open HTTP Port in Security Group

- Go to EC2 Console → select your instance → click on its Security Group
- Click Edit Inbound Rules → Add Rule
- Type: HTTP | Source: Anywhere IPv4 → click Save Rules

Step 7 – Test the Web Server

- Copy the EC2 instance's Public IP → paste into a web browser
- You should see: Amazon Web Services

Step-by-Step: Windows Web Server (IIS)

- Create and connect to a Windows EC2 Machine via RDP (refer to Week 1 & 2)
- Click Start Menu → open Server Manager
- Click Add Roles and Features → click Next (multiple times)
- On the Server Roles screen → check Web Server (IIS) → click Add Features → click Next (multiple times) → click Install
- After installation, verify by browsing to localhost — the default IIS page appears

Create a Custom Web Page

- Navigate to: C:\inetpub\wwwroot
- Create a file named index.html with this content:

```
<h1>Amazon Web Services</h1>
```

Open HTTP Port in Security Group

- In EC2 Console → select the instance → click on its Security Group ID
- Click Edit Inbound Rules → Add Rule → Type: HTTP | Source: Anywhere IPv4 → Save Rules
- Copy the Public IP and paste in browser → your page appears

Week 4 – Add and Manage EBS Volumes (Linux & Windows)

Pre-Lab Questions & Answers

Q1. What is Elastic Block Store (EBS)?
Amazon Elastic Block Store (EBS) is a block-level storage service for use with EC2 instances. It provides persistent storage that can be attached to instances, similar to adding a hard drive. A single EC2 machine can have multiple EBS volumes attached.
Q2. How do you attach and detach an additional volume on a Linux Server?
Attach: In the EC2 Console, select the volume → Actions → Attach Volume → choose the instance. Then SSH into the instance, format the volume (mkfs), create a directory, and mount it (sudo mount /dev/xvdb dirname). Detach: Unmount first (sudo umount dirname), then in the EC2 Console → select the volume → Actions → Detach Volume.
Q3. What is a Repository?
A Linux repository is a centralized storage location from which the OS retrieves and installs software packages and updates.
Q4. What is a YUM Repository?
YUM repositories are warehouses of Linux software. They allow installing packages not available by default in Linux. Use: yum install <packagename>
Q5. How to Create Empty Files in Linux?
Use the touch command: sudo touch file1 file2 file3 — This creates multiple empty files at once.

Step-by-Step: Add 5GB EBS Volume to Linux Server

A. Create EC2 Instance with an Additional Volume

- During EC2 launch → scroll to Configure Storage → click Add New Volume → set size to 5 GB → check Delete on Termination → launch the instance
- In the EC2 Console → go to Volumes to verify both volumes appear

B. Mount and Use the Volume

- Connect to the Linux EC2 Machine via PuTTY
- Check current disk usage:

```
df -hT
```

- List available block devices:

```
lsblk
```

Note the name of the new EBS volume (e.g., xvdb)

- Get information about the device:

```
sudo file -s /dev/xvdb
```

- Create a file system on the volume:

```
sudo mkfs -t ext4 /dev/xvdb
```

- Create a mount directory:

```
mkdir test
```

- Mount the volume:

```
sudo mount /dev/xvdb test
```

- Navigate into the directory and create files:

```
cd test
sudo touch file1 file2 file3
ls
df -hT
```

C. Detach and Re-Attach Volume to Another Instance

- Navigate back and unmount:

```
cd
sudo umount test
```

- In EC2 Console → Volumes → select the volume → Actions → Detach Volume → confirm
- Create a second Linux EC2 instance in the same Availability Zone
- Select the volume → Actions → Attach Volume → choose the new instance → click Attach
- Connect to the second instance via PuTTY and verify:

```
df -hT
lsblk
```

Note the volume name (e.g., xvdf)

```
sudo file -s /dev/xvdf
mkdir training
sudo mount /dev/xvdf training
cd training
ls
```

You will see file1, file2, file3 — the data persists from the first instance!

Step-by-Step: Add 5GB EBS Volume to Windows Server

A. Create Windows EC2 with Additional Volume

- During EC2 launch → Configure Storage → Add New Volume → 5 GB → launch

B. Bring Volume Online and Format

- Connect via RDP → open Server Manager → click File and Storage Services → click Disks
- Select the new disk → right-click → Bring Online → click Yes
- Right-click the disk → Initialize → click Yes

- Right-click the disk → New Volume → click Next multiple times → click Create
- The new drive (e.g., D:) appears in This PC → open it and create a file

C. Take Volume Offline and Detach

- Server Manager → File and Storage Services → Disks → right-click the EBS disk → Take Offline → click Yes
- EC2 Console → select the volume → Actions → Detach Volume

D. Attach Volume to a Second Windows Instance

- Create a new Windows EC2 instance in the same Availability Zone
- Select the volume → Actions → Attach Volume → choose the new instance → click Attach
- Connect via RDP → Server Manager → File and Storage Services → Disks
- The disk shows as Offline → right-click → Bring Online → click Yes
- Open This PC → D: drive → your file from the first instance is present

Week 5 – IAM: Users, Groups, and Policies

Pre-Lab Questions & Answers

Q1. What is Identity and Access Management (IAM)?
AWS IAM enables you to manage access to AWS services and resources securely. You can create and manage AWS users, groups, and roles, and use permissions to allow or deny access to AWS resources. IAM also keeps account credentials private.
Q2. What are the key features of AWS IAM?
Shared Account Access, Multi-Factor Authentication (MFA), Identity Federation (Facebook/Google login), Free to use (no additional charge), Password Policy management, and Granular Permissions (each user gets only the permissions they need).
Q3. Define AWS Users, Groups, and Roles.
User: Represents a person or service interacting with AWS. Starts with no permissions. Groups: A collection of IAM users; permissions assigned to the group apply to all members. Roles: A set of permissions that can be assumed by users, applications, or services — not tied to a specific user.
Q4. What is MFA in AWS IAM?
Multi-Factor Authentication (MFA) adds a second layer of security. After entering username and password, the user must also enter a one-time password (OTP) generated by their phone or MFA device. This prevents unauthorized access even if credentials are stolen.
Q5. What are IAM Users' Access and Secret Keys?
Access keys allow programmatic access to AWS via CLI or SDK. They consist of an Access Key ID (like a username) and a Secret Access Key (like a password). They are used instead of username/password for API and CLI access.

Key IAM Concepts

Root User vs IAM User

- Root User: Created when you first sign up for AWS. Has full unrestricted access to all services including billing and security credentials. Should be protected and rarely used.
- IAM User: Created by the root user or admin. Starts with no permissions. Access is granted per job function. Cannot be renamed from the AWS Console (must use CLI/SDK).

Step-by-Step: Create IAM Users with Different Policies

Goal: Create two IAM users — one with EC2 Full Access, one with S3 Full Access.

Create User 1 (EC2 Full Access)

- Go to IAM → click Users → click Add Users
- Enter a username (e.g., ec2-user-1)
- Select AWS Access Type → check Password – AWS Management Console Access
- Select Custom Password → enter a password → uncheck Require Password Reset
- Click Next: Permissions → select Attach Existing Policies Directly
- Search for AmazonEC2FullAccess → select it
- Click Next: Tags → Next: Review → Create User

Create User 2 (S3 Full Access)

- Repeat the steps above but search for and attach AmazonS3FullAccess policy

Test IAM Login

- Note down the Account ID (Console ID) from your root account
- Go to the AWS login page → select IAM User → enter Account ID → enter IAM username and password → click Sign In
- Change the password when prompted
- Verify: User 1 can access EC2 only; User 2 can access S3 only

---

Step-by-Step: Create IAM Groups and Attach Policies

Key facts about IAM Groups: A group can contain many users; a user can belong to multiple groups. Groups cannot be nested. No default group includes all users. Default limit is 300 groups (max 500).

Create a Group

- Go to IAM → click User Groups → click Create Group
- Enter a group name
- Attach a Permission Policy → search AmazonS3FullAccess → select it → click Create Group

Create a New User and Add to Group

- Go to Users → Add User → enter username and password
- Click Next: Permissions → select Add User to Group → select the group → click Next → Create User
- Log in as the user and verify S3 access

Move User Between Groups

- Create a second group → attach AmazonEC2ReadOnlyAccess policy
- Open the first group → select the user → click Remove Users
- Open the second group → click Add Users → select the user → click Add Users
- Log in and verify the user now has EC2 read-only access

Week 6 – Amazon S3: Create, Configure, and Versioning

Pre-Lab Questions & Answers

Q6. What is the full form of S3?
S3 stands for Simple Storage Service. It is an object-based cloud storage service provided by AWS.
Q7. List out the features of S3.
Global service, low latency and high throughput, safe object-based storage, 99.999999999% durability (11 nines), max object size 5 TB, unlimited total storage per bucket, globally unique bucket names, lifecycle management, versioning, encryption, and access control via ACLs and bucket policies.
Q8. What are the naming conventions for creating a bucket?
No uppercase letters, no underscores, 3–63 characters long, cannot be an IP address, must start with a lowercase letter or number.
Q9. What is Versioning?
Versioning in Amazon S3 is a feature that keeps multiple variants of an object in the same bucket. It preserves, retrieves, and restores every version of every object. It is enabled at the bucket level.
Q10. What are the advantages of Versioning?
1. You can recover deleted objects. 2. You can maintain and roll back to different versions of a file.

Key S3 Concepts

- S3 is a global service; buckets are defined at the region level
- Max single object size: 5 TB (use multipart upload for objects > 5 GB)
- Default bucket limit: 100 per account (can be raised to 1,000 via service quota request)
- Bucket names are globally unique and cannot be changed after creation
- Bucket ownership cannot be transferred to another account

Step-by-Step: Create S3 Bucket

- Go to S3 → click Create Bucket
- Enter a globally unique bucket name (follow naming conventions)
- Select the AWS Region
- Select ACLs Enabled

- Uncheck Block All Public Access → check the Acknowledge checkbox
- Click Create Bucket

Step-by-Step: Upload an Object

- Click the bucket name → click Upload → click Add Files → select your file
- Click Permissions → select Grant Public Read Access → click Upload

---

Step-by-Step: View and Change Object Permissions

- Click the object name → go to the Properties tab → copy the Object URL → paste in browser (public access works)
- To restrict access: click the object → Permissions tab → click Edit → uncheck Read Access for public → click Save Changes
- Pasting the Object URL now returns an Access Denied error

Step-by-Step: Enable and Use Versioning

Enable Versioning

- Open your bucket → go to Properties tab → click Edit for Bucket Versioning → select Enable → click Save Changes

Make Object Public (required before versioning demo)

- Open bucket → Permissions tab → Edit Block Public Access → uncheck Block All Public Access → Save → type confirm → Confirm
- Click the object → Permissions tab → Edit → check Read checkbox for Public → Save Changes → verify Object URL works in browser

Recover a Deleted Object

- Open bucket → click Show Versions → select the file → click Delete → type delete → click Delete Objects
- After deletion, click Show Versions → you will see a Delete Marker file
- Select the Delete Marker → click Delete → type permanently delete → click Delete Objects
- The original file is restored and visible again

Maintain Different Versions of a File

- Edit your file locally (add more lines) → upload the new version to the same bucket with the same file name
- Grant Public Read Access during upload
- Click Show Versions → you will see both the old and new versions listed

- To recover the old version: select it → click Download → re-upload if needed

Week 7 – S3 Lifecycle Management

Pre-Lab Questions & Answers

Q1. What is the maximum size of an object that can be stored in an S3 bucket?
The maximum size for a single object is 5 terabytes (5 TB). For objects larger than 5 GB, you should use multipart upload. The total volume of data a bucket can hold is unlimited.
Q2. What are the different Storage Classes in Amazon S3?
S3 Standard (frequent access), S3 Intelligent-Tiering (auto-tiering), S3 Standard-IA (infrequent access), S3 One Zone-IA (lower cost, single AZ), S3 Glacier Instant Retrieval, S3 Glacier Flexible Retrieval, and S3 Glacier Deep Archive (lowest cost, long-term archive).
Q3. How can you secure data in S3?
Use Bucket Policies (resource-based JSON policies), Access Control Lists (ACLs), IAM policies for users/roles, S3 Block Public Access settings, and Server-Side Encryption (SSE) with S3-managed keys, KMS keys, or customer-provided keys.
Q4. What tools does AWS offer for monitoring S3?
Amazon CloudWatch (metrics and alarms), AWS CloudTrail (API activity logs), S3 Server Access Logging, S3 Storage Lens (organization-wide visibility), and AWS Config (configuration compliance tracking).
Q5. What are ACLs in AWS S3?
Access Control Lists (ACLs) are a legacy access control mechanism in S3. They define which AWS accounts or groups have access to buckets and objects, and what permissions they have (read, write, full control). ACLs must be enabled to use them in newer bucket configurations.

Lifecycle Management – Concept

Lifecycle Management is a way of efficiently managing S3 objects to minimize storage costs over their lifetime. Up to 1,000 lifecycle rules can be configured per bucket.

There are two types of lifecycle actions:

- Transition Actions: Move objects to a cheaper storage class after a set number of days (e.g., move from S3 Standard to S3 Glacier after 30 days)
- Expiration Actions: Permanently delete objects after a set number of days

---

Step-by-Step: Apply a Lifecycle Rule to an S3 Bucket

Step 1 – Create the Bucket and Upload Object

- Go to S3 → Create Bucket
- Select ACLs Enabled → Give Public Access → Create Bucket
- Upload an object → Grant Public Read Access

Step 2 – Create the Lifecycle Rule

- Open your bucket → click the Management tab
- Click Create Lifecycle Rule
- Enter a Rule Name
- Select: This rule applies to all objects in the bucket → check the Acknowledge checkbox
- Select the desired rule actions based on your requirements:
  - Transition to another storage class: select the target storage class (e.g., S3 Glacier) and enter the number of days after which the transition occurs
  - Expire current versions: enter the number of days after which objects are permanently deleted
- Review the Timeline Summary to confirm the rule schedule
- Click Create Rule

Note: Always review cost implications before setting lifecycle rules, as transitions to certain storage classes may have minimum storage duration charges.

Week 11 – Create and Update a DynamoDB Table

Pre-Lab Questions & Answers

Q1. What is DynamoDB?
Amazon DynamoDB is a fast, fully-managed NoSQL database service provided by AWS. It makes it simple and cost-effective to store and retrieve any amount of data and serve any level of request traffic. It requires no server management and scales automatically.

Q2. What is the purpose of DynamoDB?
DynamoDB is used to store and retrieve any amount of structured data quickly and reliably. It is designed for applications that need consistent, single-digit millisecond latency at any scale. Common use cases include web/mobile apps, gaming, IoT, real-time analytics, and session management.

Q3. What type of data can be stored in DynamoDB?
DynamoDB stores semi-structured and unstructured NoSQL data. Supported data types include: Scalar Types (String, Number, Boolean, Binary, Null), Document Types (List, Map), and Set Types (String Set, Number Set, Binary Set). Each item (row) can have a different set of attributes.

Q4. What is the Partition Key in DynamoDB?
The Partition Key (also called the Primary Key) is a mandatory attribute that uniquely identifies each item in a DynamoDB table. DynamoDB uses its value as input to an internal hash function, which determines the physical partition where the item is stored. Each value must be unique across all items in the table when no Sort Key is used.

Q5. What is the Sort Key in DynamoDB?
The Sort Key (also called the Range Key) is an optional second attribute used together with the Partition Key to form a Composite Primary Key. When a Sort Key is defined, multiple items can share the same Partition Key as long as their Sort Key values are different. Items with the same Partition Key are stored together and sorted by the Sort Key value.

Q6. What are the features of DynamoDB?
Fully managed NoSQL database (no server setup needed), single-digit millisecond response times at any scale, automatic scaling of throughput and storage, built-in security

with IAM integration and encryption at rest, supports two key constraint types: Primary Key (Partition Key) and Composite Primary Key (Partition Key + Sort Key), supports backup and restore, global tables for multi-region replication, DynamoDB Streams for change data capture, and a free tier available.

Key Concepts

Amazon DynamoDB is a fast, fully-managed NoSQL database service. It is serverless — there is no infrastructure to manage. DynamoDB automatically scales to handle any amount of traffic and stores data across multiple availability zones for high availability.

DynamoDB Key Constraints

- **Primary Key (Partition Key)** — A single attribute that uniquely identifies every item in the table. No two items can share the same Partition Key value.
- **Composite Primary Key (Partition Key + Sort Key)** — Two attributes together form the unique identifier. Multiple items can share the same Partition Key as long as their Sort Key values differ. Items with the same Partition Key are sorted by Sort Key.

Step-by-Step: Create a DynamoDB Table

A. Create the Table

- Go to the AWS Console → search for and open DynamoDB
- Click on Create Table
- Enter the Table Name (e.g., StudentRecords)
- Enter the Partition Key name (e.g., StudentID) and select its data type (String, Number, or Binary)
- Optionally enter a Sort Key name (e.g., CourseName) and select its data type
- Leave Table Settings as Default Settings
- Click on Create Table — wait for the table status to become Active

B. Add Items (Rows) to the Table

- In the DynamoDB Console, click on the Table Name to open it
- Click on Explore Table Items
- Click on Create Item
- Enter the value for the Partition Key (and Sort Key if applicable) — these are mandatory
- To add more columns (attributes) to this item, click on Add New Attribute
- Select the attribute type (e.g., String, Number, Boolean) — then enter the Attribute Name and its Value
- Repeat Add New Attribute for each additional column needed
- Click on Create Item to save it

- Repeat the above steps to create more items (each item can have a different set of attributes)

C. Delete the DynamoDB Table

- Go to DynamoDB → Tables
- Select the checkbox next to the table you want to delete
- Click on Delete → confirm deletion by typing the table name when prompted → click Delete

Note: Deleting a table permanently removes all its data and cannot be undone. Make sure to back up any important data before deleting a table.